

## Part 1

A.

In this analysis, I will be using a summation of the following variables: HighBlood, Stroke, Overweight, Arthritis, Diabetes, Hyperlipidemia, BackPain, Anxiety, Allergic\_rhinitis, Reflux\_esophagitis, and Asthma into a new variable called health\_score. I will also be using the variable Doc\_visits. My goal is to answer the following question: is there a positive correlation between Doc\_visits + health\_score and readmission rates for the hospital. This way I can answer an important business question: does an increase in the number of Doc\_visits plus a high number of the health\_score column lead to a greater chance of readmission rates. In other words, is a high number for both Doc\_visits and health\_score positively correlated with readmission rates?

B.

I will be running analysis on the following variables: Doc\_visits, HighBlood, Stroke, Overweight, Arthritis, Diabetes, Hyperlipidemia, BackPain, Anxiety, Allergic\_rhinitis, Reflux\_esophagitis, and Asthma. Doc\_visits is a number while all the other data are yes or no answers except for Overweight and Anxiety, which are either 1 or 0. Since some of the data is numeric, and some is categorical, I will be re-coding all the yes or no answers into numbers 1 or 0 (1 for yes, 0 for no). This way all our independent variables will be numeric and our correlation will be consistent across all columns. We will need to create a new column that stores the “health score” from the other columns which will be called health\_score. Our analysis will therefore be looking at explaining the ReAdmis data by using Doc\_visits and health\_score.

## Part 2

C.

I will first be re-coding all the categorical data in the columns HighBlood, Stroke, Overweight, Arthritis, Diabetes, Hyperlipidemia, BackPain, Anxiety, Allergic\_rhinitis, Reflux\_esophagitis, and Asthma to numeric data with yes answers being a 1 and no answers being a 0. There are a few Python functions needed to re-code this data into the new column called health\_score, which will then have a box-plot and whisker plot to show the distribution of the data. Then I will repeat the process on Doc\_visits to show how the data is spread out. I will also be calculating Z-scores on both of the variables mentioned to identify anomalies/outliers.

Since I will be re-coding all the data from categorical to numeric and then making a new column, it would make sense to drop all rows that have a null value. I will be using Python methods to detect and drop any rows that contain null values, missing data, and altering inconsistent datatypes. Judging from the data at hand, since the data dictionary does NOT explain whether 1 or 0 represents either a yes or no answer to both the Overweight and Anxiety columns, both of these columns need to be excluded from the analysis; they will not be included in the health\_score column. Next I will be calculating Z-scores on Doc\_visits and health\_score. I will be analyzing all the data's Z-scores and dropping any individual rows that have a Z-score higher than 3 or less than -3.

Using Z-scores to determine outliers is a very common and strong tool when identifying data that doesn't meet the *general* trend. Since our goal is to find general trends within the data, using Z-scores in combination with visual plots should give us a powerful understanding of how to lower readmission rates using these two variables.

Since all the data are numeric and have no null values, I will assume that the data was recorded without error and truly represents the population in question. Having no null values makes our analysis strong since it suggests that the hospital did a good job recording the information and that the data reflects reality. The quality of the data is therefore strong and will provide us with an insightful analysis to answer the research question.

I will be using Python with the libraries of pandas, numpy, matplotlib, and scipy. These libraries give the Python language a way to interpret the data and show trends, as well as clean the data so that our analysis is consistent and without error. I will not be using R for any of the analysis since Python has all the necessary libraries included in the Anaconda package. All code is written and verified in Jupyter notebook.

Here is the code for detecting our outliers where I checked for null values:

```
df['health_score'].notnull().any()
```

```
True
```

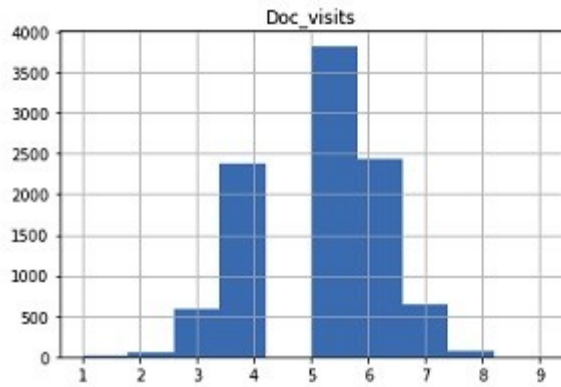
```
df['Doc_visits'].notnull().any()
```

```
True
```

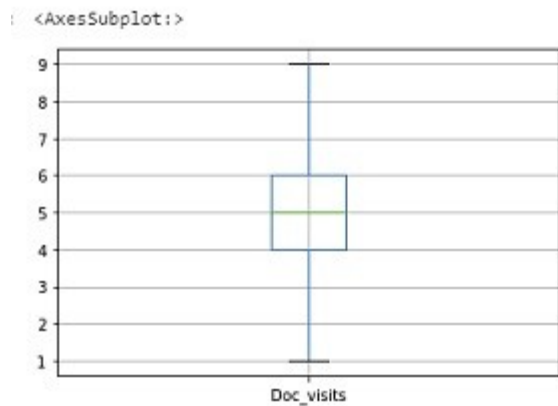
The notnull().any() function will check to see if there are any not null values (which means we have all data that is not null) and returns true if it is true.

Here are box plots and histograms for the Doc\_visits column so we can graphically represent the distribution of the data:

```
df.hist(['Doc_visits'])  
array([[<AxesSubplot:title={'center':'Doc_visits'}>]], dtype=object)
```



```
df.boxplot(['Doc_visits'])
```

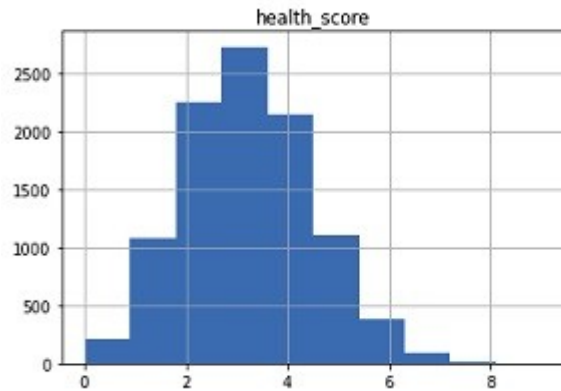


We can see from the code that our data is mostly a normal distribution with some gaps from 4-5 Doc\_visits.

Below we have a box plot and histogram for our newly made column `health_score`:

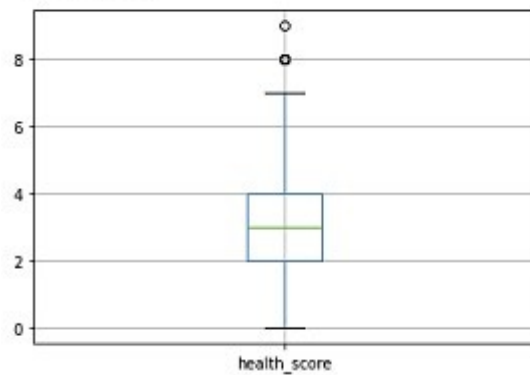
```
df.hist(['health_score'])
```

```
array([[<AxesSubplot:title={'center':'health_score'}>]], dtype=object)
```



```
df.boxplot(['health_score'])
```

```
<AxesSubplot:>
```



Again the histogram shows a basic normal distribution implying that most of our sample population is less than one standard deviation from the mean.

### Part 3

D.

Firstly, I dropped both the Anxiety column and the Overweight column from the analysis; they weren't included at all. Then I check for null values using the `isnull().any()` function which returns true if the columns have no null values; there were no null values or missing data found in the data set.

Next, I re-code all categorical data into numeric data and then sum them to create our new column called `health_score`. I then show how this data is distributed with a box plot and histogram and move onto the next stage which is calculating Z-scores and dropping outliers.

In order to calculate Z-scores I had to make a new column that stores Z-scores. To get a new Z-score column for `health_score`, we have to create an array out of the `health_score` column, save it to a new variable, and then run the `stats.zscore` method on the new array. We then save the output to a new column and add it to our data frame.

We repeat the same process that we did for `health_score` but now for `Doc_visits` and get a new column. The columns are titled `zscore_hs` (stores Z-scores for `health_score`) and `zscore_dv` (stores Z-Scores for `Doc_visits`).

One problem I encountered was that, after we had created our new column to store Z-scores, the datatype for the new columns were floats and not ints. I changed the datatype back to int using the `astype()` method and now all our columns are ready for the next stage.

Next I run the `drop()` method and remove any rows that have a Z-score higher than the absolute value of 3, which resulted in only one row being removed and this is verified by running the `info()` method which showed that we still have 9,999 rows of data remaining. This is good news since our data has now been fully cleaned and all data outliers/anomalies have been removed.

Below is where I re-code all categorical data to numeric data and create our new `health_score` column:

```
df = df.replace(to_replace={'HighBlood': {'Yes': 1, 'No': 0}}, value=None)
print(df['HighBlood'])
```

```
***
```

```
df = df.replace(to_replace={'Stroke': {'Yes': 1, 'No': 0}}, value=None)
print(df['Stroke'])
```

```
***
```

```
df = df.replace(to_replace={'Arthritis': {'Yes': 1, 'No': 0}}, value=None)
print(df['Arthritis'])
```

```
***
```

```
df = df.replace(to_replace={'Diabetes': {'Yes': 1, 'No': 0}}, value=None)
print(df['Diabetes'])
```

```
***
```

```
df = df.replace(to_replace={'Hyperlipidemia': {'Yes': 1, 'No': 0}}, value=None)
print(df['Hyperlipidemia'])
```

```
***
```

```
df = df.replace(to_replace={'BackPain': {'Yes': 1, 'No': 0}}, value=None)
print(df['BackPain'])
```

```
***
```

```
df = df.replace(to_replace={'Allergic_rhinitis': {'Yes': 1, 'No': 0}}, value=None)
print(df['Allergic_rhinitis'])
```

```
***
```

```
df = df.replace(to_replace={'Reflux_esophagitis': {'Yes': 1, 'No': 0}}, value=None)
print(df['Reflux_esophagitis'])
```

```
***
```

```
df = df.replace(to_replace={'Asthma': {'Yes': 1, 'No': 0}}, value=None)
print(df['Asthma'])
```

```
***
```

Code for creating the health\_score column:

```
df['health_score'] = df.apply(lambda x: x['HighBlood'] + x['Stroke'] + x['Arthritis'] + x['Diabetes'] + x['Hyperlipidemia'] + x['BackPain'] + x['Allergic_rhinitis'] + x['Reflux_esophagitis'] + x['Asthma'], axis=1)
print(df['health_score'])
```

0	6
1	2
2	2
3	4
4	2
...	...
9995	2
9996	4
9997	2
9998	1
9999	3

Name: health\_score, Length: 10000, dtype: int64

Next I create a Z-score column for both Doc\_visits and health\_score:

```
import numpy as np
import scipy.stats as stats
```

```
hs_col_4z=df.loc[ : , 'health_score']
```

```
print(hs_col_4z)
```

```
***
```

```
df['zscore_hs']=stats.zscore(hs_col_4z)
df['zscore_hs'].head()
```

```
***
```

```
dv_col_4z=df.loc[ : , 'Doc_visits']
```

```
print(dv_col_4z)
```

```
***
```

```
df['zscore_dv']=stats.zscore(dv_col_4z)
df['zscore_dv'].head()
```



Next I had to convert both the new Z-score columns from a float to an int to make our datatypes consistent and then filter for any int values that were more than the absolute value of 3 and drop those rows:

```
df['zscore_dv'] = df['zscore_dv'].astype(int)
```

```
df.info()
```

```
***
```

```
df['zscore_hs'] = df['zscore_hs'].astype(int)
df.info()
```

```
***
```

```
df.drop(df[df['zscore_hs'] > 3].index, inplace=True)
df.info()
```

```
***
```

```
df.drop(df[df['zscore_hs'] < -3].index, inplace=True)
df.info()
```

```
***
```

```
df.drop(df[df['zscore_dv'] > 3].index, inplace=True)
df.info()
```

```
***
```

```
df.drop(df[df['zscore_dv'] < -3].index, inplace=True)
df.info()
```


```
***
```

I run `df.info()` after each change so that I can verify that the values were changed and dropped correctly.

Below are the histogram charts for our Z-score columns so that we can better understand how the data is distributed:


```
df.hist(['zscore_hs'])
```

```
array([[<AxesSubplot:title={'center':'zscore_hs'}>]], dtype=object)
```



```
df.hist(['zscore_dv'])
```

```
array([[<AxesSubplot:title={'center':'zscore_dv'}>]], dtype=object)
```



In summary, we first took our raw data and excluded two columns completely (Anxiety and Overweight) because the data dictionary did not indicate whether or not a 1 or 0 represented a yes or a no answer. Then we filtered for any missing data or null values for the remaining columns and found none.

We then took our remaining columns (HighBlood, Stroke, Arthritis, Diabetes, Hyperlipidemia, BackPain, Reflux\_esophagitis, and Asthma) and transformed all 'Yes' rows to be the number 1 and all 'No' rows to be number 0. We then summed all those columns into a new column labeled health\_score so that we could run a Z-score function on the data and identify outliers.

Next we created a new column to represent the Z-scores for all the data from health\_score and Doc\_visits then dropped all rows for both those columns in which the Z-score was higher than the absolute value of 3, resulting in only 1 row out of 10,000 being removed.

Some of the limitations to this method of cleaning data is that using, or relying entirely on, Z-scores can be problematic when we run further analysis on the data. Z-scores assume a normal distribution of our variables which may not always be the case. The limitations of using a Z-score approach to removing outliers is that it removes some of the extreme cases from the analysis which may be helpful in discovering the causal relationship between health\_score + Doc\_visits and readmission rates. Although Z-scores aren't perfect, we only had to remove 1 row of data in this analysis which I think gives strength to our summary.

Below is the cleaned data set: it has no null values, it has consistent datatypes, and all outliers removed according to Z-score (1 row in total).

```
df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 9999 entries, 0 to 9999
Data columns (total 56 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Unnamed: 0            9999 non-null   int64
 1   CaseOrder             9999 non-null   int64
 2   Customer_id           9999 non-null   object
 3   Interaction            9999 non-null   object
 4   UID                   9999 non-null   object
 5   City                  9999 non-null   object
 6   State                 9999 non-null   object
 7   County                9999 non-null   object
 8   Zip                   9999 non-null   int64
 9   Lat                   9999 non-null   float64
10  Lng                   9999 non-null   float64
11  Population            9999 non-null   int64
12  Area                  9999 non-null   object
13  Timezone              9999 non-null   object
14  Job                   9999 non-null   object
15  Children              7411 non-null   float64
16  Age                   7585 non-null   float64
17  Education             9999 non-null   object
18  Employment            9999 non-null   object
19  Income                7536 non-null   float64
20  Marital               9999 non-null   object
21  Gender                9999 non-null   object
22  ReAdmis               9999 non-null   object
23  VitD_levels           9999 non-null   float64
24  Doc_visits            9999 non-null   int64
25  Full_meals_eaten      9999 non-null   int64
26  VitD_supp             9999 non-null   int64
27  Soft_drink            7533 non-null   object
28  Initial_admin         9999 non-null   object
29  HighBlood             9999 non-null   int64
30  Stroke                9999 non-null   int64
31  Complication_risk     9999 non-null   object
32  Overweight            9017 non-null   float64
33  Arthritis             9999 non-null   int64
34  Diabetes              9999 non-null   int64
35  Hyperlipidemia        9999 non-null   int64
36  BackPain              9999 non-null   int64
37  Anxiety               9016 non-null   float64
38  Allergic_rhinitis     9999 non-null   int64
39  Reflux_esophagitis    9999 non-null   int64
40  Asthma                9999 non-null   int64
41  Services              9999 non-null   object
42  Initial_days          8943 non-null   float64
43  TotalCharge           9999 non-null   float64
44  Additional_charges    9999 non-null   float64
45  Item1                 9999 non-null   int64
46  Item2                 9999 non-null   int64
47  Item3                 9999 non-null   int64
48  Item4                 9999 non-null   int64
49  Item5                 9999 non-null   int64
50  Item6                 9999 non-null   int64
51  Item7                 9999 non-null   int64
52  Item8                 9999 non-null   int64
53  health_score          9999 non-null   int64
54  zscore_hs             9999 non-null   int32
55  zscore_dv             9999 non-null   int32
dtypes: float64(11), int32(2), int64(25), object(18)
memory usage: 4.5 MB
```

In conclusion, the result of this cleaning has led to two columns: health\_score and Doc\_visits which have been cleaned of any missing values, null values, and outliers/anomalies. Our data is now cleaned and ready for the next step in the data analytics life cycle.

E.

In this section I am going to be doing something similar to what I did before by grouping data but I will be doing the grouping according to PCA. The components in the data set are going to be the column I created labeled health\_score, as well as Doc\_visits, and VitD\_levels. Using 3 components total we can then analyze the data using these variables and see if we can learn something about readmissions using PCA.


PCA has to use integers or floating ints, so we are somewhat limited in which variables we can choose, I chose these variables because I thought they would be the best in explaining and answering our research question. The PCA shows us that PC1 has the most weight on explaining the variance on Doc\_visits while PC3 has the most weight on explaining our other two columns (health\_score and VitD\_levels). The reason we should use PCA is the it can sometimes better explain the trends we see in data and help simplify our analysis. We can see from our scree plot that we really only need 1 component to explain most of the variance in the data because our other principal components have eigenvalues under 1.

This analysis may not benefit a ton from using PCA because a lot of our data is categorical and not numeric. While we can estimate some numerical values out of our categorical data (like I did for health\_score) it may not always be the most comprehensive approach for this data set. PCA works great

when you have a lot of numeric variables that can be logically grouped together which makes it a strong tool for performing data analytics.


Below is the code for and the graphs showing the PCA:

```
plt.plot(pca.explained_variance_ratio_)
plt.xlabel('# of components')
plt.ylabel('explained variance')
plt.show()
```



```
cov_matrix = np.dot(newdf_normalized.T, newdf_normalized) / newdf.shape[0]
eigenvalues = [np.dot(eigenvector.T, np.dot(cov_matrix, eigenvector)) for eigenvector in pca.components_]
```

```
plt.plot(eigenvalues)
plt.xlabel('number of components')
plt.ylabel('eigenvalues')
plt.show()
```



```
loadings = pd.DataFrame(pca.components_.T,
                        columns=['pc1', 'pc2', 'pc3'],
                        index=newdf.columns)
loadings
```

	pc1	pc2	pc3
Doc_visits	0.412976	-0.856269	0.310248
health_score	-0.683552	-0.066306	0.726884
VitD_levels	0.601837	0.512256	0.612687

Part 4.

G. No third party code was used other than that contained within the WGU lessons or methods found on Python panda docs and stack overflow:

[https://pandas.pydata.org/docs/user\\_guide/index.html](https://pandas.pydata.org/docs/user_guide/index.html)

<https://stackoverflow.com/>

H. No other works cited.